CAPSTONE PROJECT REPORT

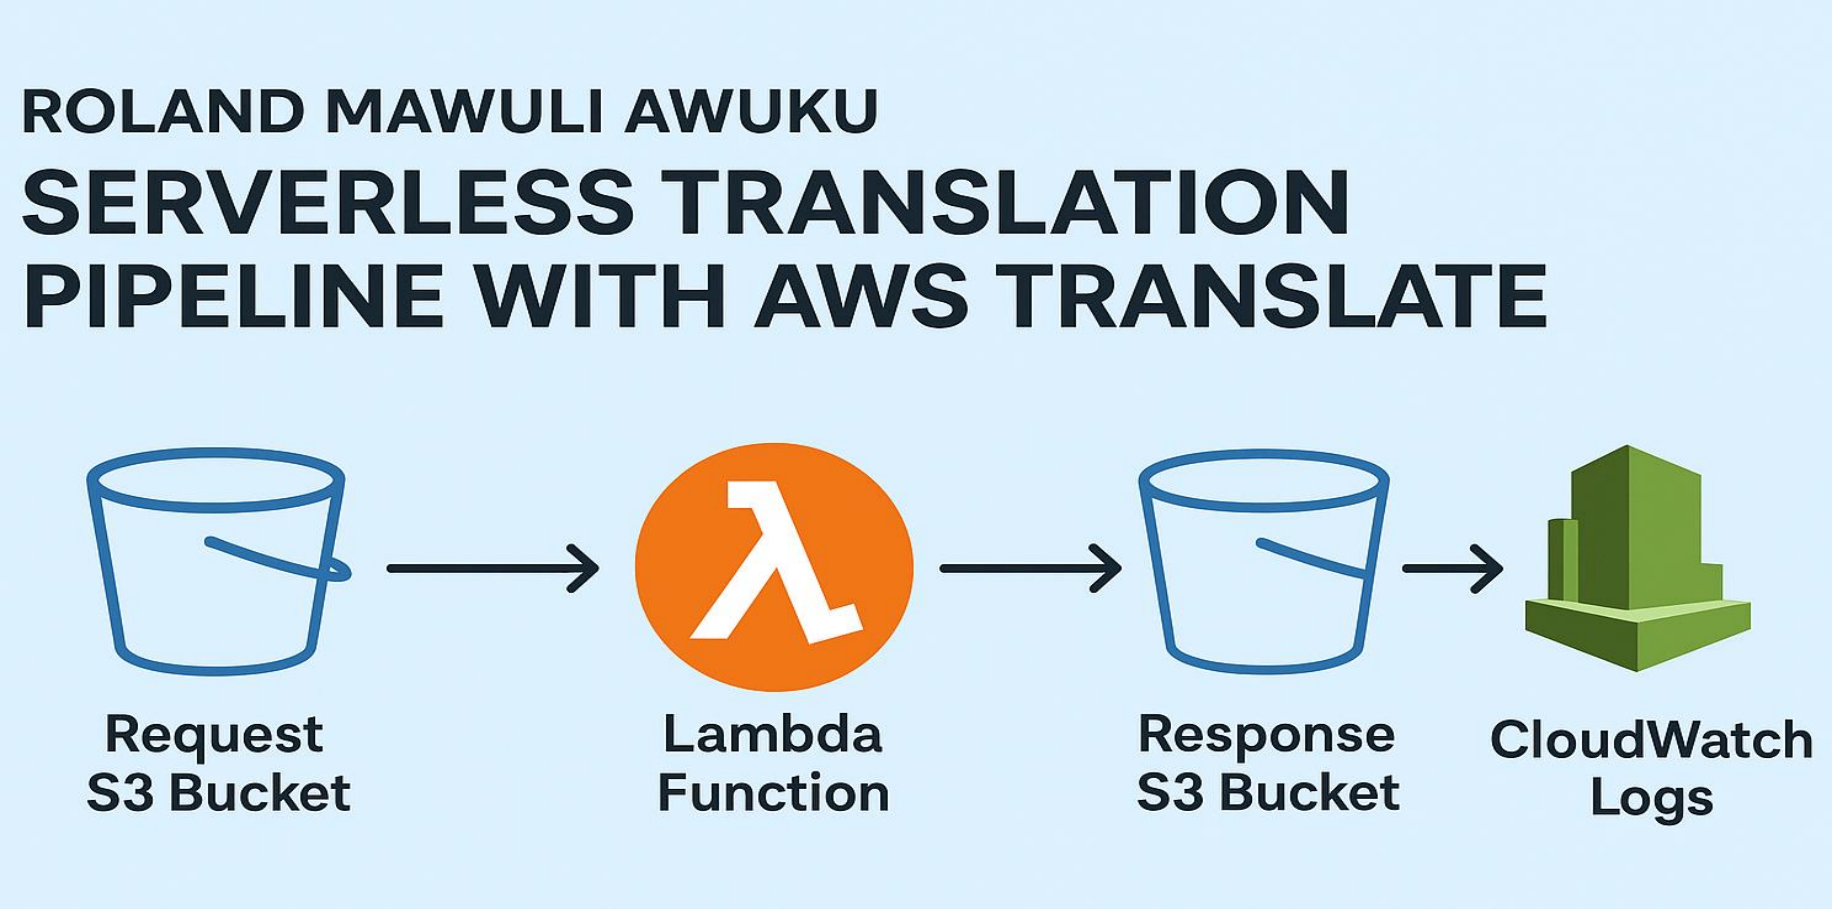

- The goal is to architect, build, and deploy an Infrastructure-as-Code (IaC) solution on AWS that enables automated, real-time language translation.
- This implementation leverages Amazon S3, AWS Lambda, and Amazon Translate, with all resources provisioned and managed through Terraform.
- The solution is entirely serverless and optimized to stay within AWS Free Tier limits when handling small-scale test datasets.

- Automate the translation of text contained in uploaded JSON files.
- Save both the original requests and their translated outputs securely in Amazon S3.
- Leverage a serverless architecture to ensure scalability while minimizing operational overhead.
- Provision and deploy infrastructure consistently and reproducibly using Terraform (IaC).
- Apply best practices for IAM, CloudWatch logging, and cost optimization throughout the solution.

The solution is built as an event-driven pipeline:

- Request S3 Bucket. Holds uploaded JSON files containing translation requests.
- S3 Event Notification. Automatically triggers the Lambda function upon file upload.
- Lambda Function (Python, Boto3). Processes the JSON file, invokes Amazon Translate, and generates translated output.
- Response S3 Bucket. Stores the resulting translated JSON files.
- CloudWatch Logs. Records execution details, errors, and debugging information.
- Terraform. Manages infrastructure as code, provisioning S3 buckets, IAM roles, and the Lambda function


PHASE 1:
ENVIRONMENT
PREPARATION
& AWS SETUP

The first stage established the groundwork for the translation pipeline by configuring core AWS services, access controls, and the development environment.

Researched Core AWS Services

Amazon Translate (real-time neural machine translation): Reviewed supported languages, character limits (5,000 bytes per request), and Free Tier allowance (2M characters/month for 12 months).

Amazon S3 (scalable object storage): Studied key features such as versioning, encryption, bucket policies, lifecycle management, and event notifications to support later automation steps.

Installed and Configured Essential Tools

AWS CLI - enabled local interaction with AWS resources.

Terraform - for Infrastructure-as-Code resource provisioning.

Python - selected for Lambda function development.

AWS Credentials Setup

Configured local authentication using `aws configure`

```
PS C:\Users\user\Downloads\CapProject\CapProject> aws configure
AWS Access Key ID [****************25JV]: AKIASL6G3UHFBHV24XU7
AWS Secret Access Key [****************zdNZ]: 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
Default region name [us-east-1]: eu-north-1
Default output format [JSON]:
PS C:\Users\user\Downloads\CapProject\CapProject> terraform init
```

Drafted IAM Policy with Scoped Permissions

`translate:TranslateText` - access to Amazon Translate.

`s3:GetObject, s3:PutObject, s3:ListBucket` - permissions for Amazon S3.

CloudWatch log permissions - to capture Lambda execution details.

PHASE 2: INFRASTRUCTURE AS-CODE (IAC) DESIGN

In this phase, Terraform was used to fully automate the provisioning of AWS resources, including S3 buckets, IAM roles, and the Lambda deployment. This ensured a reproducible and error-free setup compared to manual configuration.

ACTIVITIES:

Created Terraform provider configuration

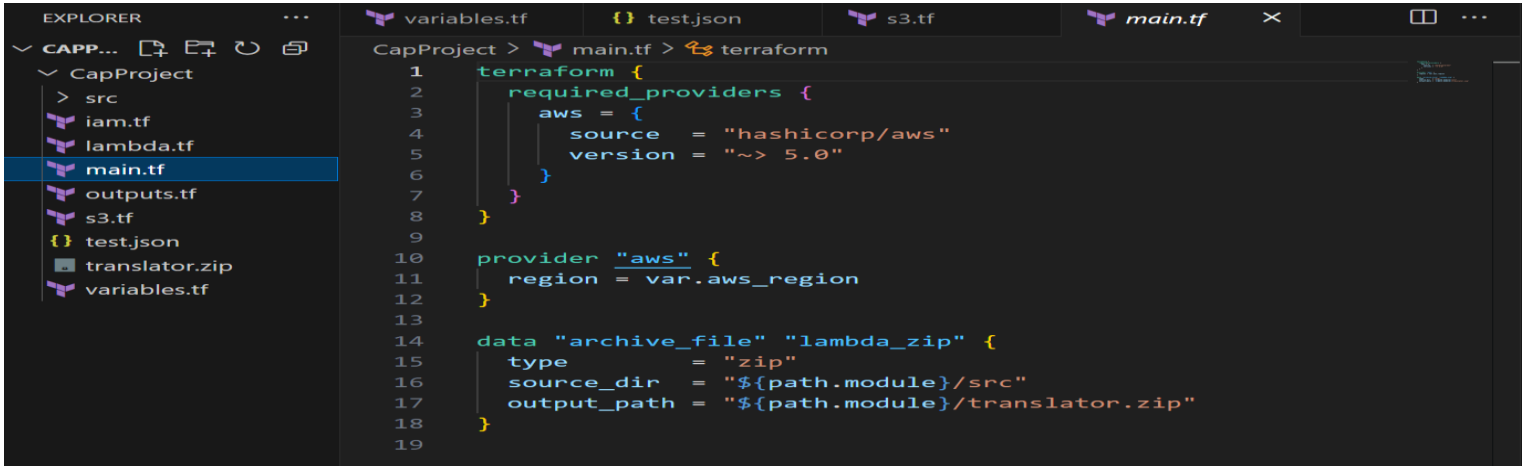

Defined variables for region, input/output buckets, and Lambda function:

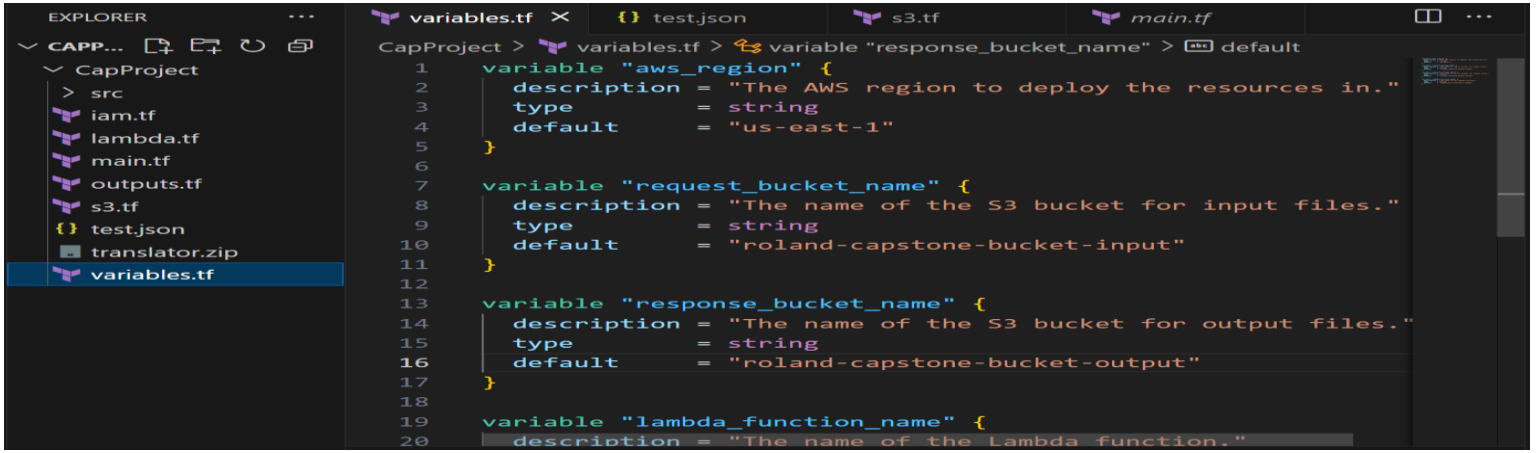

S3 Buckets with Lifecycle Policies

Two Amazon S3 buckets were provisioned using Terraform:

request-bucket. For storing the incoming JSON files requiring translation.

response-bucket. For storing the translated JSON files after processing.

Lifecycle rules were added to automatically transition or expire files after a set period, ensuring cost optimization within Free Tier limits.

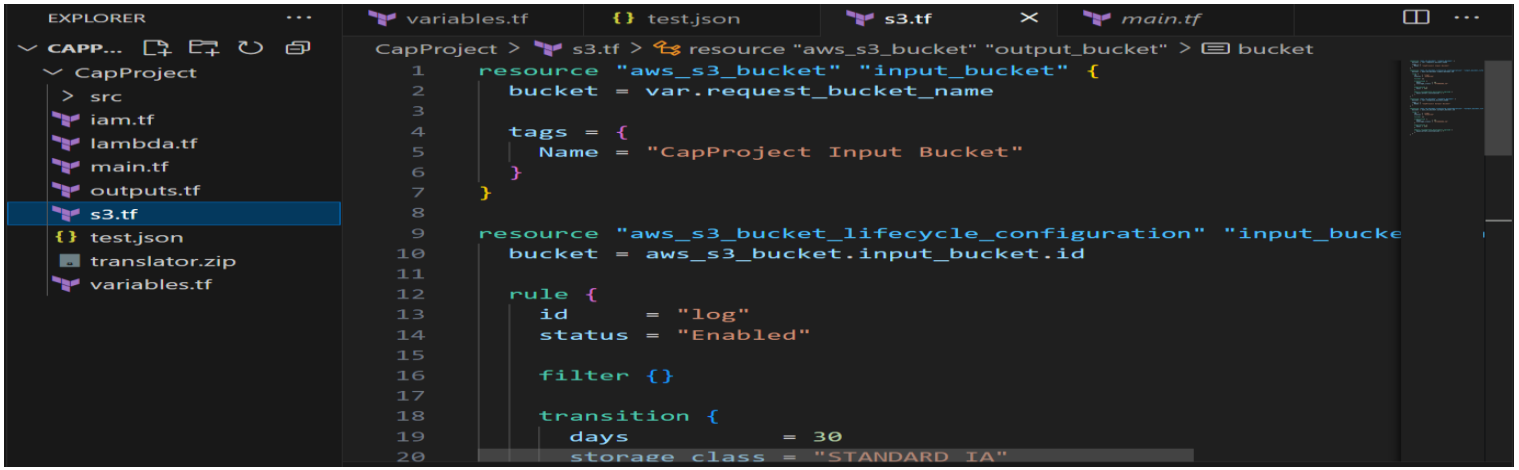

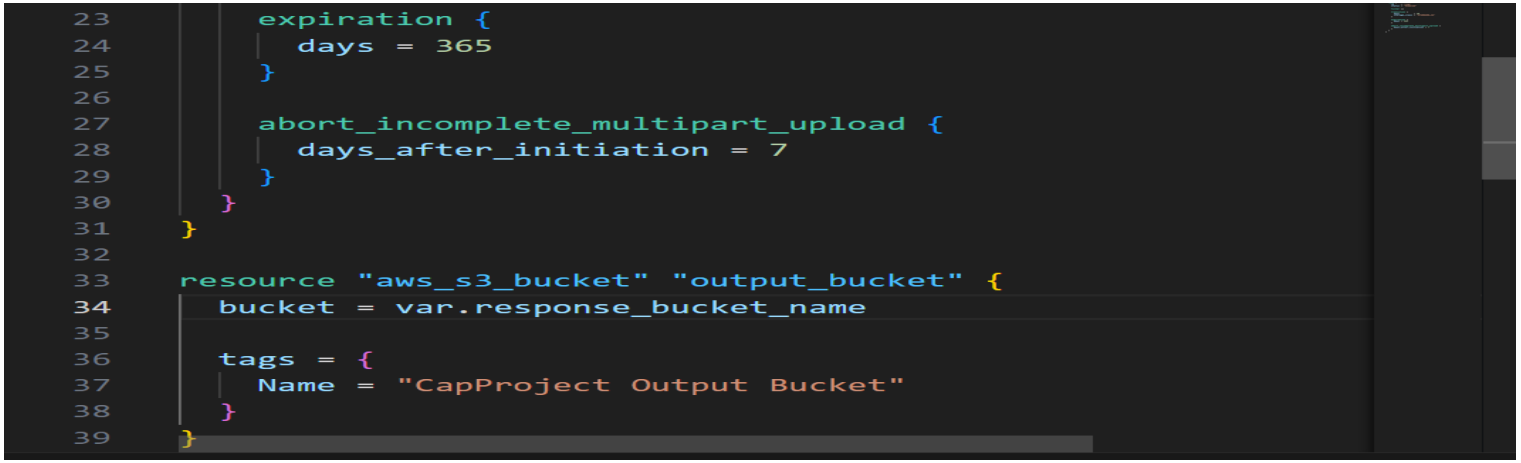

```
41  resource "aws_s3_bucket_lifecycle_configuration" "output_buck
42    bucket = aws_s3_bucket.output_bucket.id
43
44    rule {
45      id      = "log"
46      status = "Enabled"
47
48      filter {}
49
50      transition {
51        days          = 30
52        storage_class = "STANDARD_IA"
53      }
54
55      expiration {
56        days = 365
57      }
58
59      abort_incomplete_multipart_upload {
```

```
57        }
58
59        abort_incomplete_multipart_upload {
60          days_after_initiation = 7
61        }
62      }
63    }
64
```

After provisioning the S3 buckets, Terraform outputs the names of the input and output buckets:

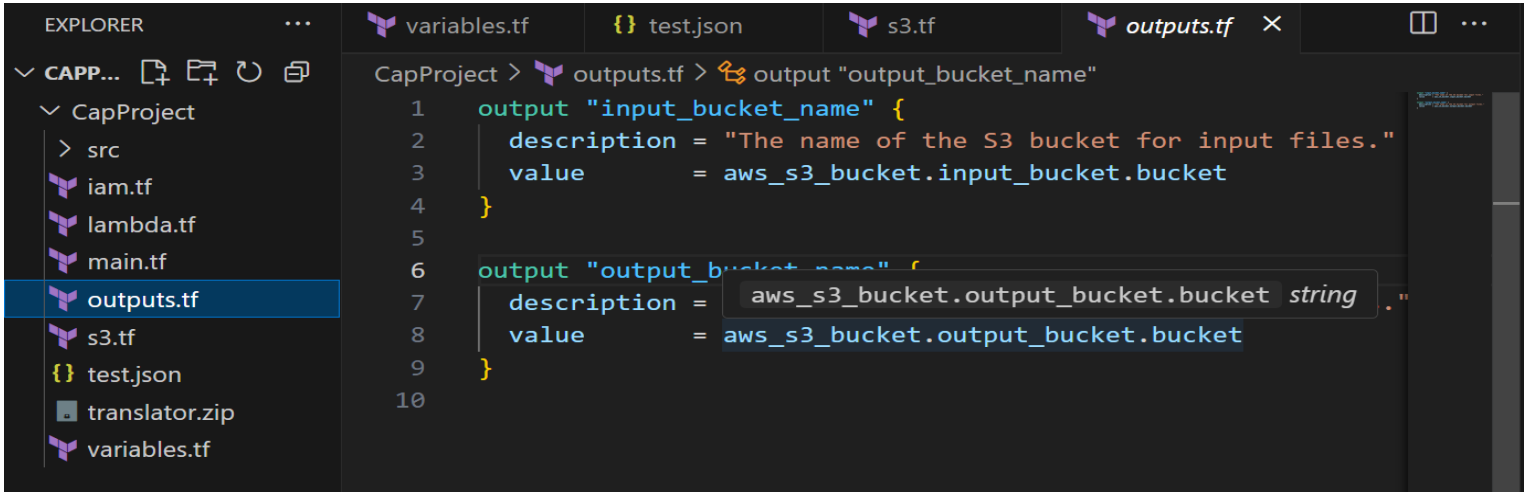

IAM Role and Policy for Lambda Execution

Using Terraform, a least-privilege IAM role and policy were provisioned for the Lambda function. This guaranteed that the function operated with only the permissions it needed.

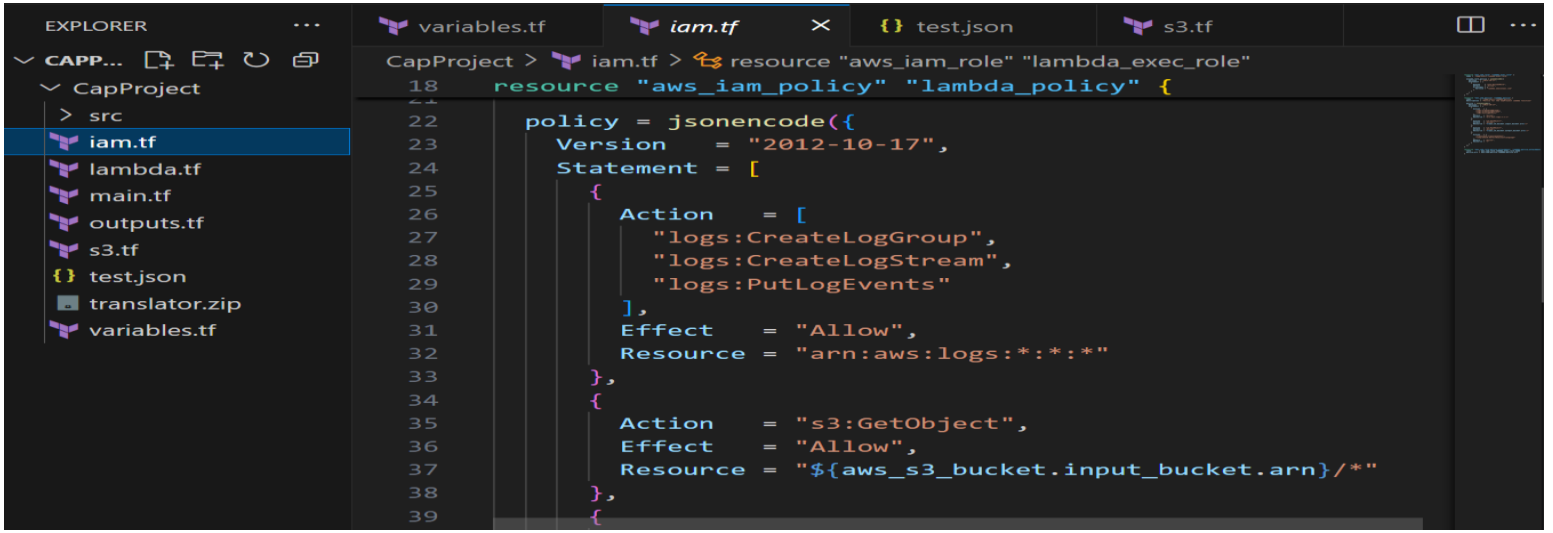
EXPLORER
CAPP...
CapProject
src
iam.tf
lambda.tf
main.tf
outputs.tf
s3.tf
test.json
translator.zip
variables.tf
variables.tf
iam.tf
test.json
s3.tf
CapProject > iam.tf > resource "aws_iam_role" "lambda_exec_role"
18 resource "aws_iam_policy" "lambda_policy" {
22 policy = jsonencode({
23 Version = "2012-10-17",
24 Statement = [
25 {
26 Action = [
27 "logs:CreateLogGroup",
28 "logs:CreateLogStream",
29 "logs:PutLogEvents"
30],
31 Effect = "Allow",
32 Resource = "arn:aws:logs:*:*:*"
33 },
34 {
35 Action = "s3:GetObject",
36 Effect = "Allow",
37 Resource = "${aws_s3_bucket.input_bucket.arn}/*"
38 },
39 {


EXPLORER
CAPP...
CapProject
src
iam.tf
lambda.tf
main.tf
outputs.tf
s3.tf
test.json
translator.zip
variables.tf
variables.tf
iam.tf
test.json
s3.tf
CapProject > iam.tf > resource "aws_iam_role" "lambda_exec_role"
18 resource "aws_iam_policy" "lambda_policy" {
22 policy = jsonencode({
43 },
44 {
45 Action = [
46 "translate:TranslateText",
47 "comprehend:DetectDominantLanguage"
48],
49 Effect = "Allow",
50 Resource = "*"
51 }
52]
53 })
54 }
55
56 resource "aws_iam_role_policy_attachment" "lambda_policy_atta
57 role = aws_iam_role.lambda_exec_role.name
58 policy_arn = aws_iam_policy.lambda_policy.arn
59 }

Validation

Ran *terraform* to deploy stack.

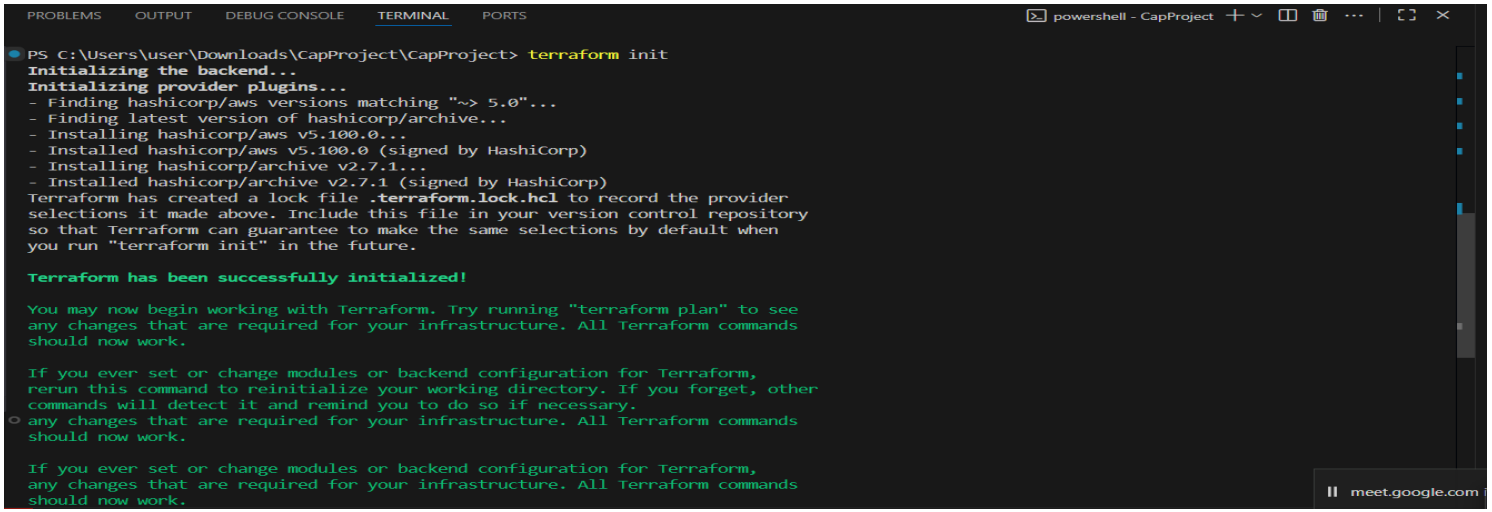

Verified resources in the AWS Console

S3 buckets

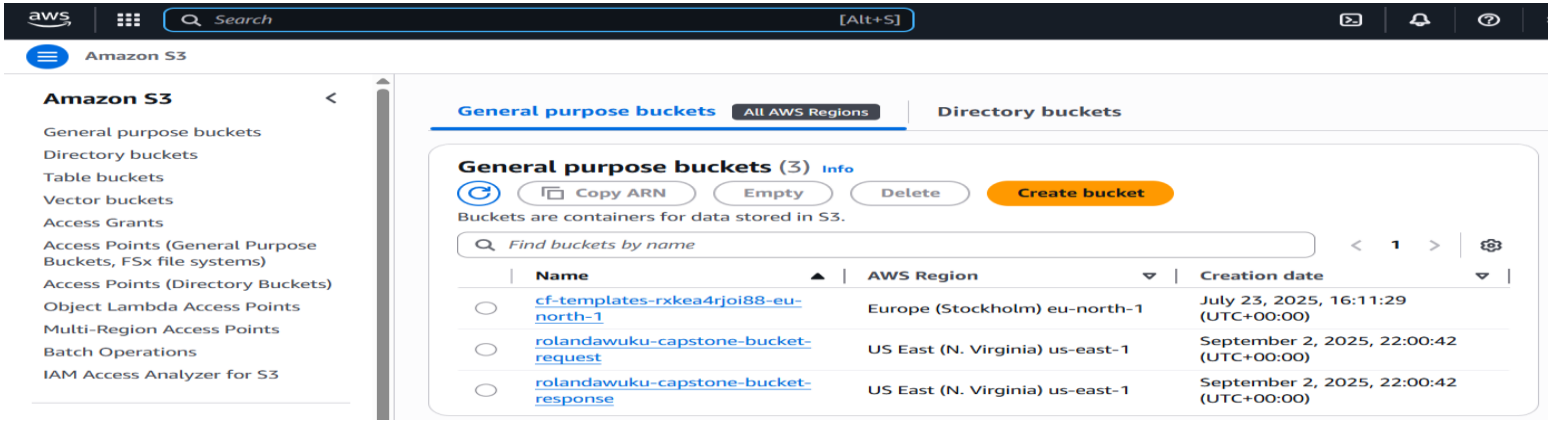

Lambda Function

Lambda execution role

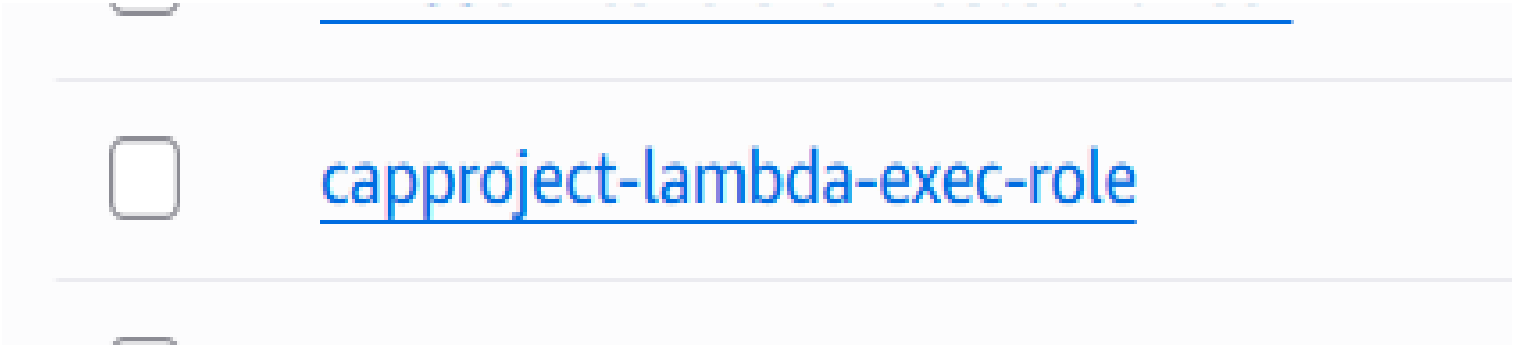

Lambda Policies

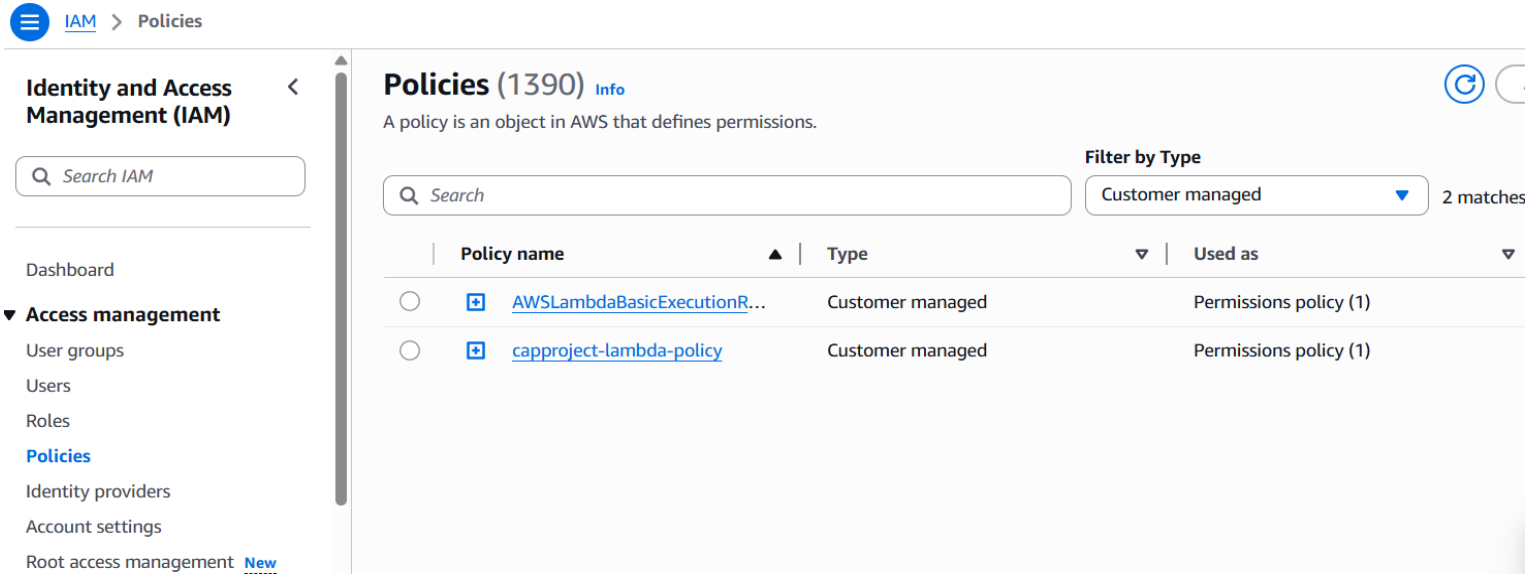

Lambda function

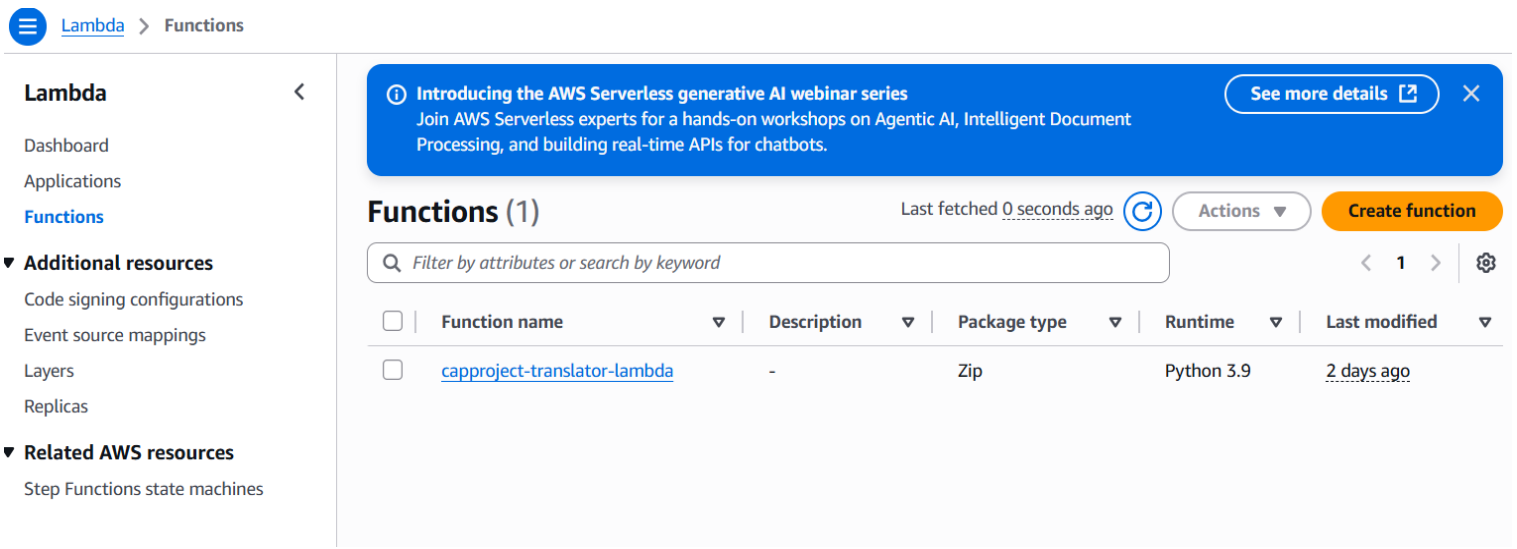

Lambda Event Logs

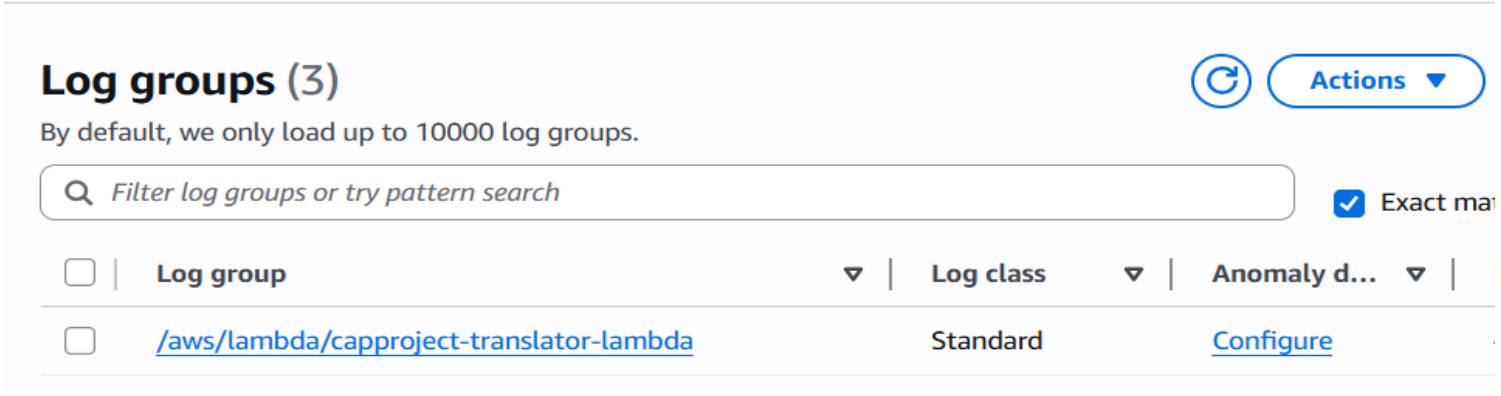

PHASE 3: BACKEND SCRIPT DEVELOPMENT (TRANSLATION LOGIC

ACTIVITIES

Developed Python script (Lambda handler) using Boto3.

Reads input JSON file with text + target language. | Uses Amazon Translate API. | Stores translated JSON in response-bucket.

JSON input:

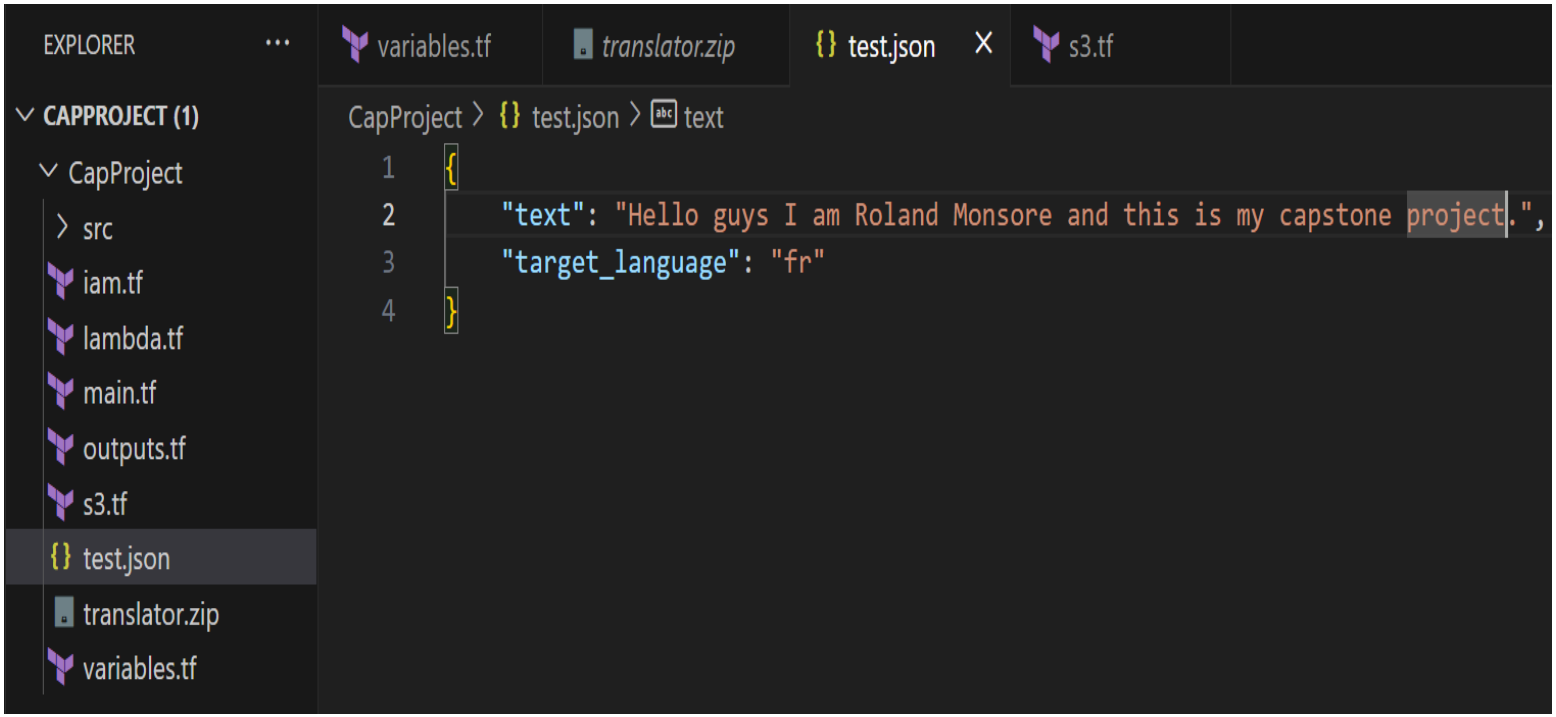

Picture of upload to request bucket:

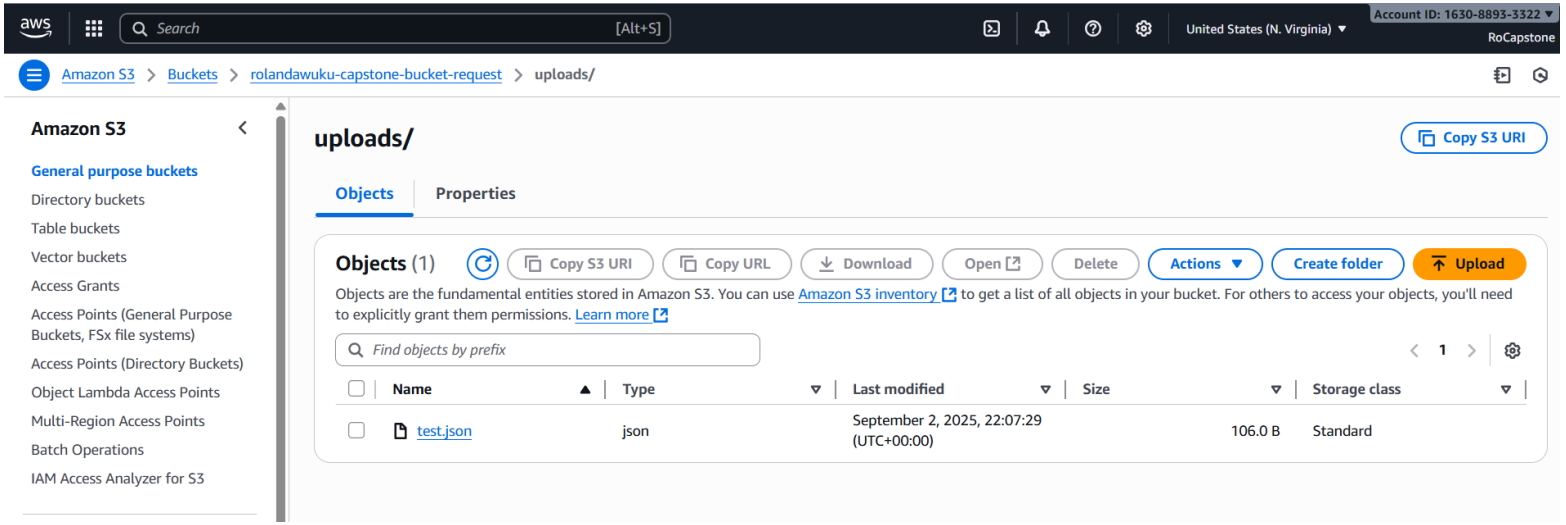

Translated JSON output file in response bucket:

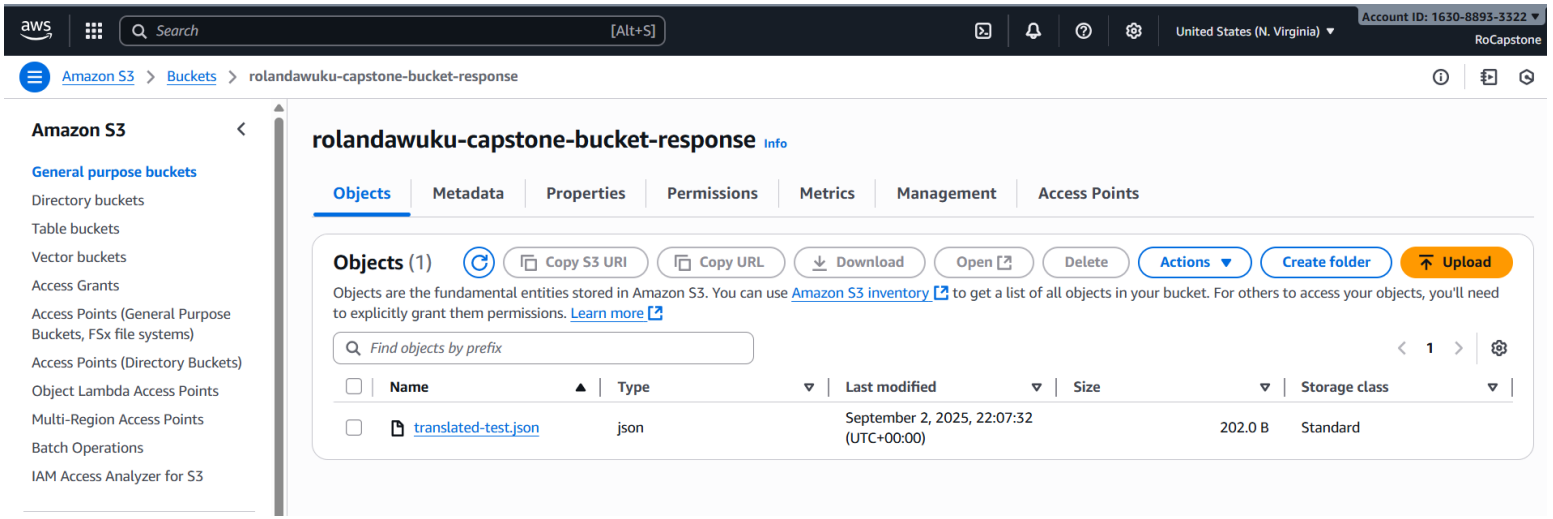

PHASE 4: AUTOMATION WITH AWS LAMBDA

Activities

Packaged the Python script using Terraform's archive_file.

Deployed AWS Lambda function with: 256 MB memory | 60-second timeout | Environment variable pointing to output bucket

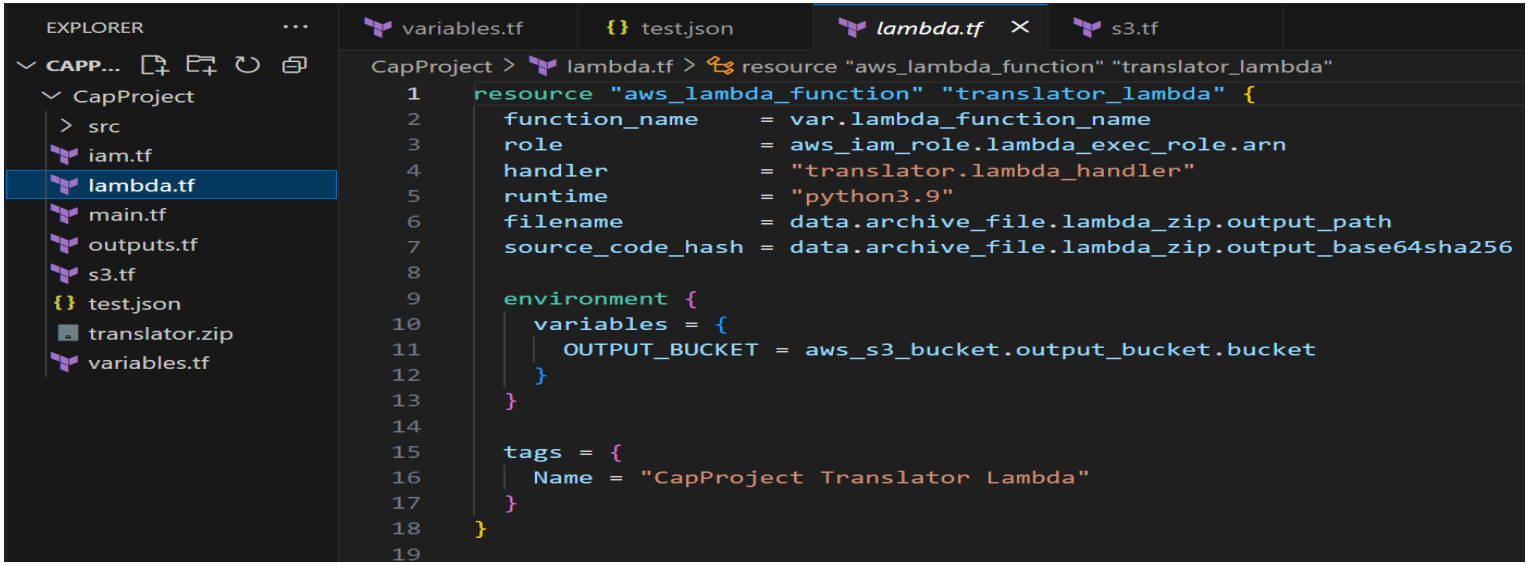

Configured S3 event notification: Whenever a file is uploaded to the input bucket, it triggers Lambda automatically.

Automated flow:

User uploads JSON → S3 triggers Lambda → Lambda (translator.py) translates → Result stored in response-bucket → CloudWatch logs capture execution details.

PHASE 5

TESTING, DEBUGGING & DOCUMENTATION

ACTIVITIES

Uploaded sample JSON files to input bucket.

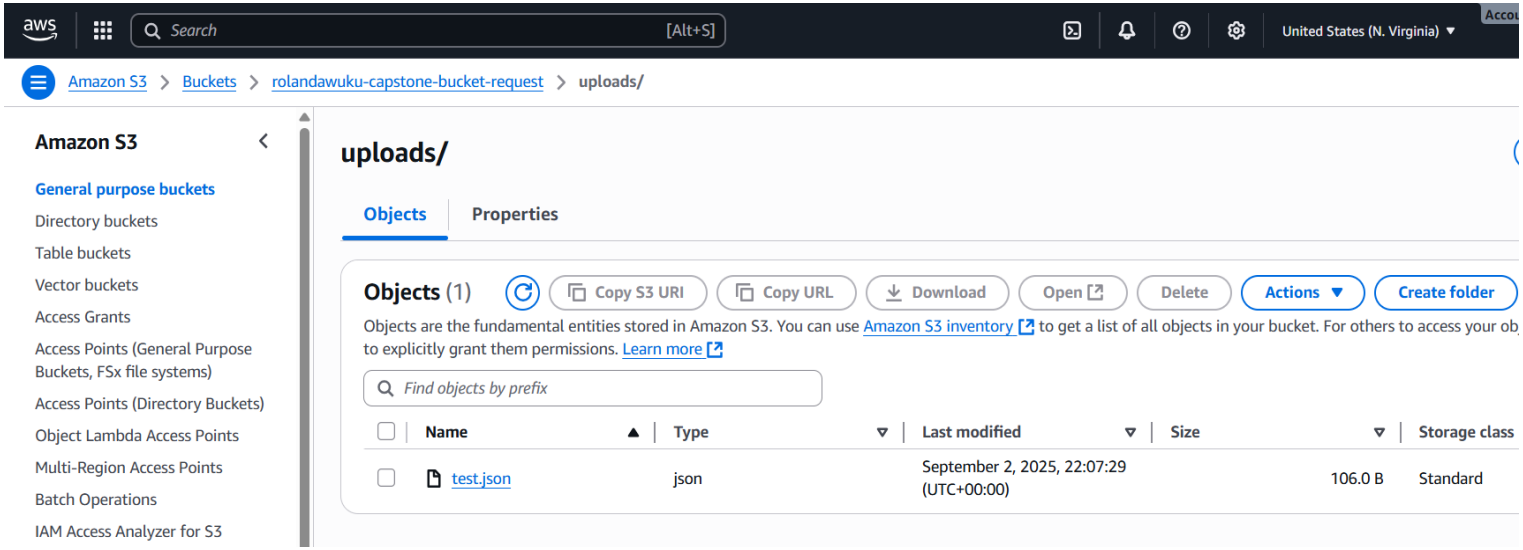

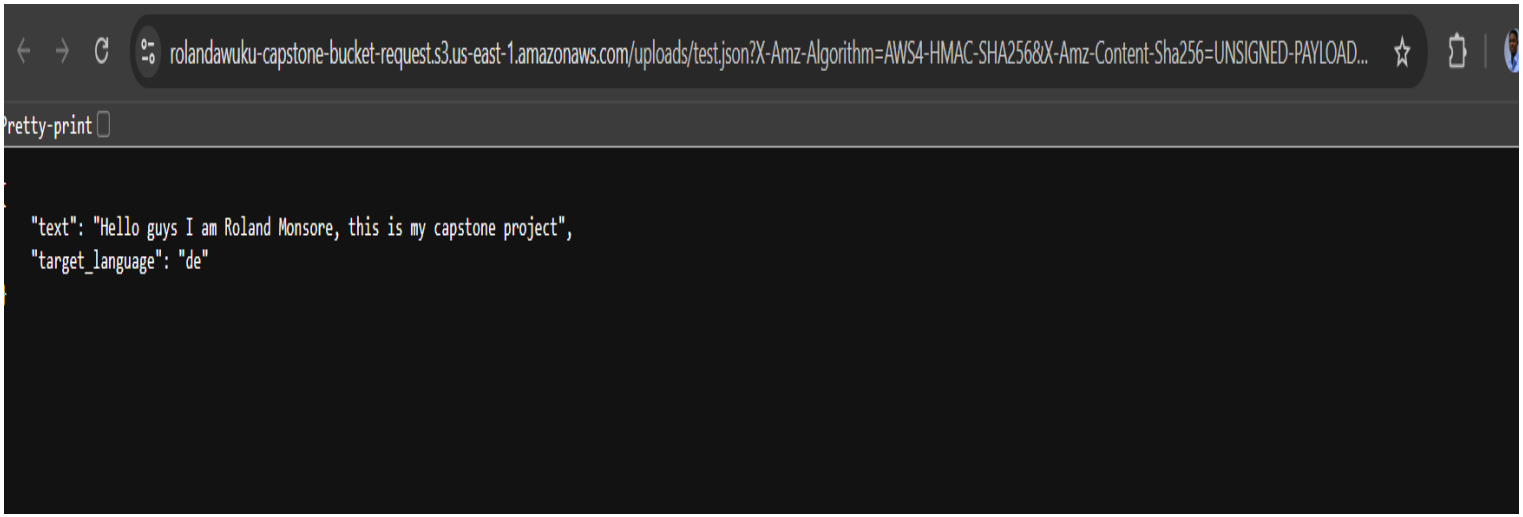

Verified translation outputs.

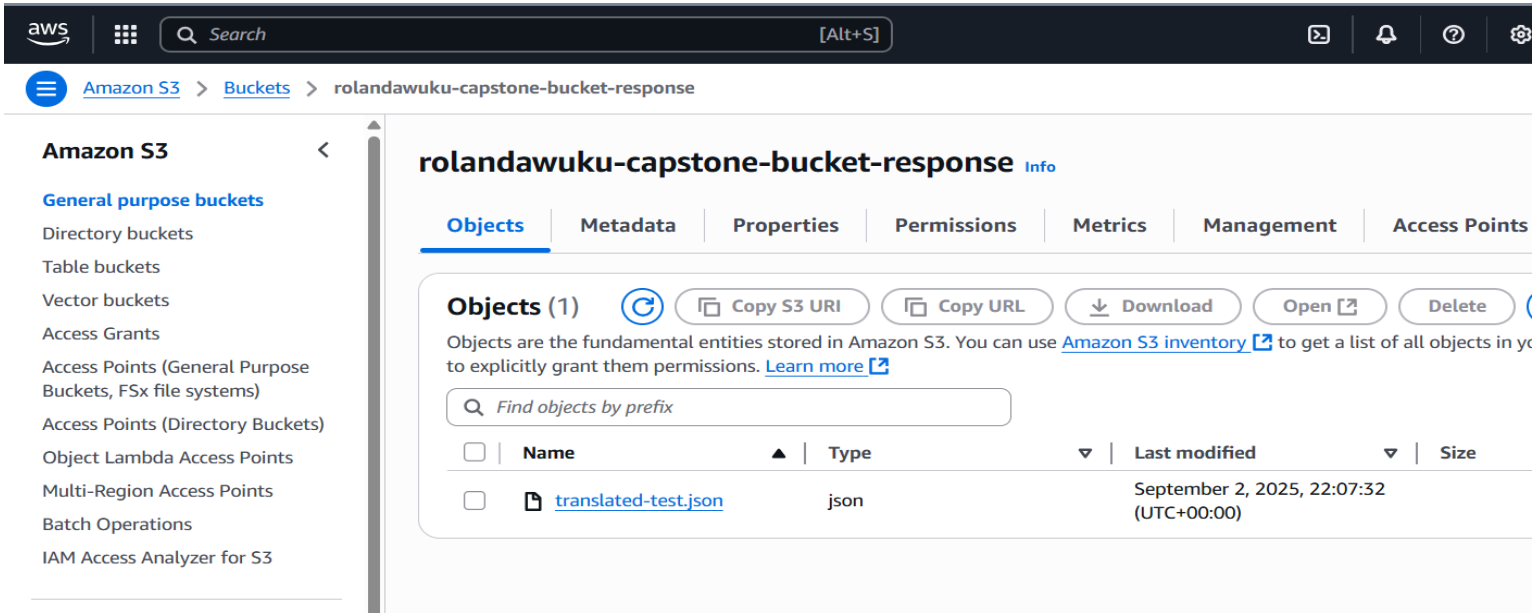

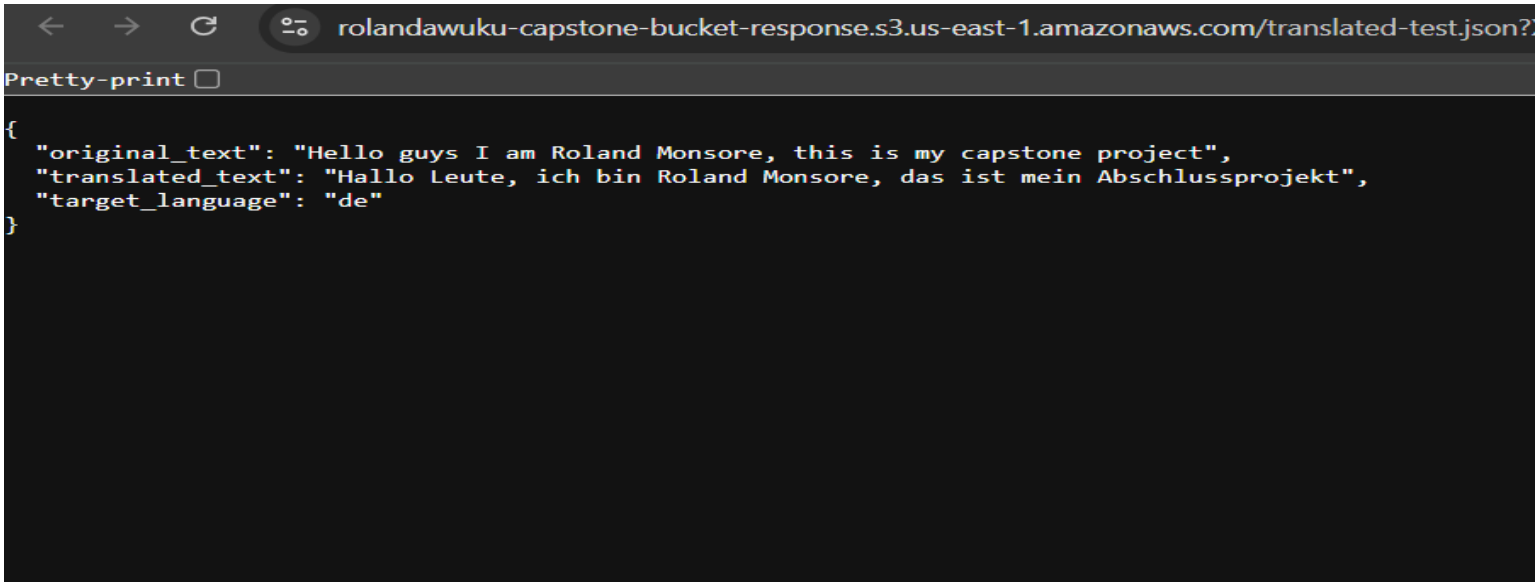

Cloudwatch logs.

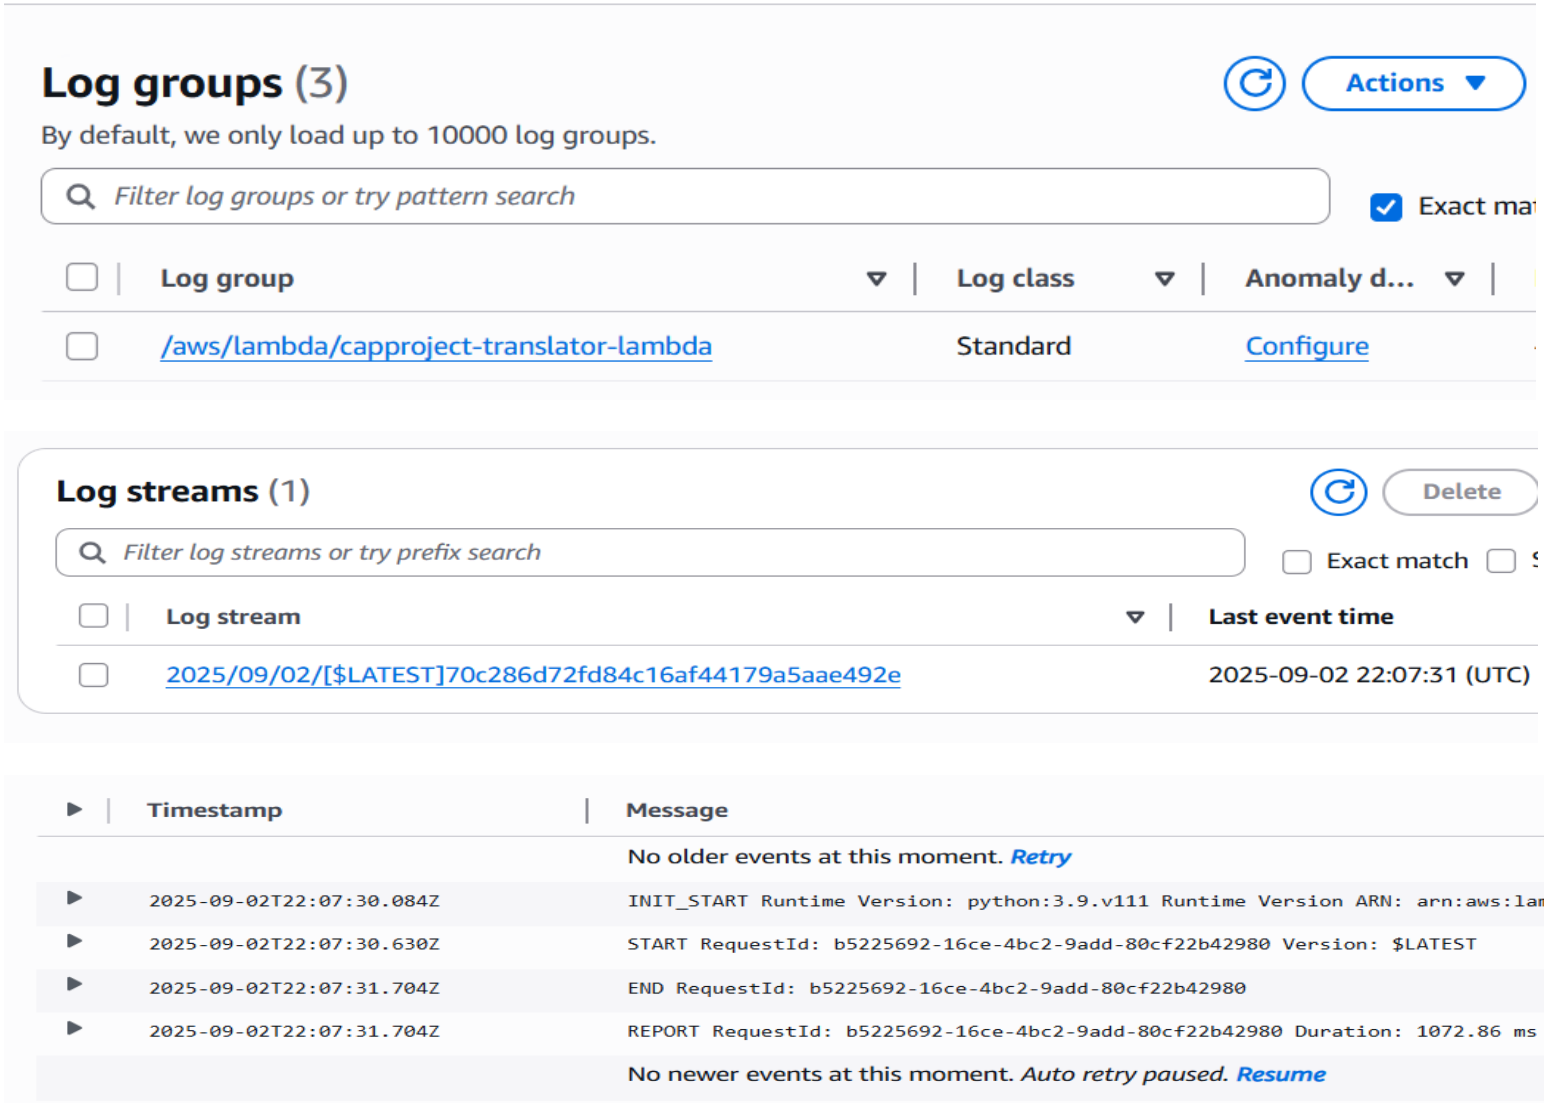

Finally ran terraform destroy to cleanup resources.

GITHUB REPOSITORY SETUP

```
PS C:\Users\user\Downloads\CapProject (1)> git init
Reinitialized existing Git repository in C:/Users/user/Downloads/CapProject (1)/.git/
PS C:\Users\user\Downloads\CapProject (1)> git add .
PS C:\Users\user\Downloads\CapProject (1)> git commit -m "first commit"
[master (root-commit) afa112f] first commit
 9 files changed, 296 insertions(+)
 create mode 100644 CapProject/iam.tf
 create mode 100644 CapProject/lambda.tf
 create mode 100644 CapProject/main.tf
 create mode 100644 CapProject/outputs.tf
 create mode 100644 CapProject/s3.tf
 create mode 100644 CapProject/src/translator.py
 create mode 100644 CapProject/test.json
 create mode 100644 CapProject/translator.zip
 create mode 100644 CapProject/main.tf
 create mode 100644 CapProject/outputs.tf
 create mode 100644 CapProject/s3.tf
 create mode 100644 CapProject/src/translator.py
```

```
PS C:\Users\user\Downloads\CapProject (1)> git branch -M main
PS C:\Users\user\Downloads\CapProject (1)> git remote add origin https://githPS C:\Users\user\Downloads\CapProje
ct (1)> git remote add origin https://github.com/R0C0d6/Ro-Cloud-Capstone.git
ub.com/R0C0d6/Ro-Cloud-Capstone.git
PS C:\Users\user\Downloads\CapProject (1)> git push -u origin main
Enumerating objects: 13, done.
Counting objects: 100% (13/13), done.
Delta compression using up to 12 threads
Compressing objects: 100% (11/11), done.
Compressing objects: 100% (11/11), done.
Writing objects: 100% (13/13), 4.30 KiB | 2.15 MiB/s, done.
Writing objects: 100% (13/13), 4.30 KiB | 2.15 MiB/s, done.
Total 13 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
Total 13 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/R0C0d6/Ro-Cloud-Capstone.git
 * [new branch]      main -> main
```

CONCLUSION

The project successfully implemented a fully serverless translation pipeline that:

- Translates JSON files uploaded to Amazon S3 into a specified target language automatically.
- Saves both the original and translated versions securely in S3.
- Utilizes Amazon Comprehend to detect the source language dynamically.
- Captures execution details and errors in CloudWatch for monitoring and troubleshooting.
- Uses Terraform for reproducible deployments, ensuring consistent infrastructure provisioning.
- Runs within AWS Free Tier limits when handling small-scale test data.
- The resulting system delivers a scalable, cost-effective, and automated solution for real-time language translation, highlighting best practices in serverless computing, Infrastructure as Code (IaC), and cloud-native design.

ARCHITECTURE DIAGRAM